



## Convolutions

reducing:

3×3 kernel  
2×2 stride  
64 output channels

non-reducing:

4×4 kernel  
1×1 stride  
16 output channels